



VIRGINIA COMMONWEALTH UNIVERSITY

Statistical analysis and modelling (SCMA 632)

A2: Multiple Regression Analysis for NSSO68

Linear Regression Analysis for IPL performance and salary

Chandhini Kerachan Muraleedharan

V01107497

Date of Submission: 23-06-2024

CONTENTS

Sl. No.	Title	Page No.
1.	Introduction	1
2.	Results	2
3.	Interpretations	2
4.	Recommendations	9
5.	Codes	10
6.	References	

1. Introduction

(i) About the data :

NSSO-Consumption Dataset: Contains numerical data on consumption patterns of various commodities (e.g., grains, oils, fruits) across Indian states and union territories, along with basic demographic information.

Ball by Ball Dataset: Provides detailed information on every ball bowled in IPL matches from 2008 to 2022, with 816 unique match IDs and 17 variables (numeric and text), including team names, runs scored, and player performance.

IPL Matches Dataset: Contains text-based information on IPL matches from 2008 to 2022, with 16 variables per match ID, detailing dates, cities, teams, toss results, and player details.

IPL Salary Dataset: Includes yearly salary information for IPL players, with columns for salaries in dollars and without the "\$" symbol, enabling analysis of salary trends across teams and years.

(ii) Objectives

NSSO Dataset: Multiple Regression Analysis

1. Perform multiple regression analysis on the NSSO68 dataset.
2. Conduct regression diagnostics to assess model assumptions and identify problems.
3. Interpret and explain the findings.
4. Address identified issues and re-evaluate the results.
5. Discuss significant changes observed after resolving issues.

IPL Data: Player Performance and Salary Analysis

1. Investigate the relationship between player performance and salary in the IPL through linear regression.
2. Conduct correlation analysis to explore the relationship between performance factors and pay.
3. Interpret the correlation results and discuss findings.
4. Provide insights into factors influencing player compensation, identify top performers and underperformers, and analyze statistical distributions for key players.

(iii) Business Significance

Multiple Regression Analysis and Regression Diagnostics provide valuable insights, accurate predictions, performance evaluation, resource optimization, informed decision-making, and risk management, enhancing operational effectiveness and business outcomes.

- **Insightful Factors**
- **Accurate Predictions**
- **Performance Evaluation**
- **Resource Optimization**
- **Informed Decision-Making**

2. Results

(i) Multiple Regression Analysis and Regression Diagnostics NSSO68

Python

```
In [8]: # Fit the regression model
X = subset_data[['hhdsz', 'Regular_salary_earner', 'MPCE_MRP', 'MPCE_URP', 'Possess_ration_card', 'Education', 'No_of_Meals_p
X = sm.add_constant(X) # Adds a constant term to the predictor
y = subset_data['foodtotal_v']

model = sm.OLS(y, X).fit()

# Print the regression results
print(model.summary())
```

OLS Regression Results

Dep. Variable:	foodtotal_v	R-squared:	0.503
Model:	OLS	Adj. R-squared:	0.502
Method:	Least Squares	F-statistic:	1302.2
Date:	Sun, 23 Jun 2024	Prob (F-statistic):	0.00
Time:	18:56:31	Log-Likelihood:	-61301.1
No. Observations:	9015	AIC:	1.228e+05
Df Residuals:	9007	BIC:	1.228e+05
Df Model:	7		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	361.0196	23.716	15.223	0.000	314.531	407.509
hhdsz	-12.9630	0.060	-15.074	0.000	-14.649	-11.277
Regular_salary_earner	-14.6088	6.197	-2.357	0.018	-26.756	-2.462
MPCE_MRP	0.0720	0.002	34.081	0.000	0.069	0.077
MPCE_URP	0.0592	0.002	30.731	0.000	0.055	0.063
Possess_ration_card	-48.0845	5.877	-8.181	0.000	-59.606	-36.563
Education	7.6343	0.638	11.965	0.000	6.384	8.885
No_of_Meals_per_day	49.8726	8.234	6.057	0.000	33.732	66.013

Omnibus: 3368.303 Durbin-Watson: 1.688
Prob(Omnibus): 0.000 Jarque-Bera (JB): 1256929.686
Skew: -0.422 Prob(JB): 0.00
Kurtosis: 60.840 Cond. No. 3.22e+04

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 3.22e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Interpretation :

The OLS regression results for `foodtotal_v` indicate that the model explains 50.3% of the variance in total food consumption. Key predictors include household size (-12.963, $p < 0.001$), indicating larger households consume less per capita; being a regular salary earner (-14.609, $p < 0.05$), which negatively affects food consumption; MPCE_MRP (0.073, $p < 0.001$) and MPCE_URP (0.059, $p < 0.001$), both showing higher expenditures increase food consumption. Possession of a ration card (-48.085, $p < 0.001$) decreases consumption, while higher education (7.634, $p < 0.001$) and more meals per day (49.873, $p < 0.001$) increase it. Diagnostic tests highlight potential normality issues (high Jarque-Bera and Kurtosis) and multicollinearity (high condition number). The model is statistically significant overall (F-statistic: 1302.2, $p < 0.001$). Further diagnostics and model adjustments are recommended to address multicollinearity and validate assumptions.

Strong multicollinearity or other numerical problems.

```
In [9]: # multicollinearity using Variance Inflation Factor (VIF)
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.columns))]
print(vif_data) # VIF Value more than 8 is problematic
```

	feature	VIF
0	const	105.477846
1	hhdsz	1.098855
2	Regular_salary_earner	1.138218
3	MPCE_MRP	2.068354
4	MPCE_URP	1.968635
5	Possess_ration_card	1.048881
6	Education	1.230296
7	No_of_Meals_per_day	1.004672

Interpretation :

A value of 1 means no inflation, while higher values suggest influence. Here, the code identifies "const" (often the baseline value in regression) with a very high VIF (105.477846), which might be less concerning on its own. It's the high VIF in other features that might truly signal problematic overlap. For instance, MPCE_MRP (2.068354) and MPCE_URP (1.968635) have concerningly high values, suggesting their effects might be heavily influenced by each other. "hhdsz" (1.098855), "Regular_salary_earner" (1.138218), "Possess_ration_card" (1.048881), "Education" (1.230296), and "No_of_Meals_per_day" (1.004672) have much lower VIF values, indicating less of a concern.

```
In [10]: # Extract the coefficients from the model
coefficients = model.params

# Construct the equation
equation = f"y = {coefficients[0]:.2f}"
for i in range(1, len(coefficients)):
    equation += f" + {coefficients[i]:.6f}*x{i}"

# Print the equation
print(equation)

y = 361.02 + -12.963010*x1 + -14.608830*x2 + 0.072781*x3 + 0.059190*x4 + -4
8.084503*x5 + 7.634267*x6 + 49.872590*x7
```

Interpretation :

- The coefficient for x5 is negative. This means that as the value of x5 increases, the value of y tends to decrease.
- The coefficients for all the other features are positive. This means that as the values of these features increase, the value of y tends to increase.

R

```
38
39 # Print the regression results
40 print(summary(model))
41
42 library(car)
43 # check for multicollinearity using variance inflation factor (VIF)
44 vif(model) # VIF value more than 5 is problematic
45
46 # Extract the coefficients from the model
47 coefficients <- coef(model)
48
49 [Top Level] >
```

R 4.3.2 - C:/Users/Chand/Downloads/Assignment1/

```
> print(summary(model))

Call:
lm(formula = foodtotal_q ~ MPCE_MRP + MPCE_URP + Age + Meals_At_Home +
    Possess_ration_card + Education, data = subset_data)

Residuals:
    Min       1Q   Median       3Q      Max
-20.802  -3.142  -0.229   2.817  32.168

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  15.6598813   1.4342413   10.919 < 2e-16 ***
MPCE_MRP      0.0020498   0.0002628    7.248 6.67e-13 ***
MPCE_URP      0.0014809   0.0003009    4.921 9.54e-07 ***
Age           0.0254118   0.0106221    2.392  0.0169 *
Meals_At_Home -0.0052875   0.0153790   -0.344  0.7310
Possess_ration_card -1.1968075  0.6525434   -1.834  0.0668 .
Education    -0.0469761   0.0601595   -0.781  0.4350
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.786 on 1529 degrees of freedom
Multiple R-squared:  0.3907,    Adjusted R-squared:  0.3883
F-statistic: 163.4 on 6 and 1529 DF,  p-value: < 2.2e-16

>
```

Interpretation :

This Regression analyzes how much people eat (foodtotal_q) based on various factors. The model considers income spent on both store-bought (MPCE_MRP) and homegrown (MPCE_URP) food, along with age, how many meals are eaten at home (Meals_At_Home), access to a government food program (Possess_ration_card), and education level. The analysis reveals that together these factors significantly influence eating habits (p-value < 2.2e-16), explaining 39% of the variation in food consumption. While age and education don't seem to have a statistically strong influence, income spent on food, eating habits, and access to the food program do.

```

49 # Construct the equation
50 equation <- paste0("y = ", round(coefficients[1], 2))
51 for (i in 2:length(coefficients)) {
52   equation <- paste0(equation, " + ", round(coefficients[i], 8), "x", 4-2)
53 }
54 # Print the equation
55 print(equation)
56
57 # Head subset data
58 head(subset_data$PC_WBP, 1)
59 head(subset_data$PC_WBP, 1)
60 head(subset_data$age, 1)
61 head(subset_data$heats_at_home, 1)
62 head(subset_data$possession_ratio_card, 1)
63 head(subset_data$education, 1)
64 head(subset_data$footcatal_q, 1)
65
66 # The model
67
68 ## R.2.3.3: Linear OLS Regression
69
70 Age      0.004438  0.000023  2.202  0.0109 *
71 heats_at_home -0.001287  0.0013790 -0.344  0.7310
72 possession_ratio_card -1.166807  0.0015434 -1.834  0.0658
73 education   -0.046976  0.0001595 -0.781  0.4350
74
75 Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
76
77 Residual standard error: 4.786 on 1229 degrees of freedom
78 Multiple R-squared:  0.3907,    Adjusted R-squared:  0.3883
79 F-statistic: 161.4 on 6 and 1229 DF, p-value: < 2.2e-16
80
81 # library(car)
82 # Check for multicollinearity using variance inflation factor (vif)
83 vif(model) # vif value more than 8 is problematic
84
85 # Extract the coefficients from the model
86 coefficients <- coef(model)
87 # Construct the equation
88 equation <- paste0("y = ", round(coefficients[1], 2))
89 for (i in 2:length(coefficients)) {
90   equation <- paste0(equation, " + ", round(coefficients[i], 8), "x", 4-2)
91 }
92 # Print the equation
93 print(equation)
94 [1] "y = 35.86 + 0.004438x2 + 0.001287x3 + 0.025412x4 + -0.001288x5 + -0.128607x6 + -0.046976x7"
95

```

Interpretation :

reveals a statistically significant linear regression model, explaining around 39% of the dependent variable's variance. While there's no multicollinearity concern based on VIF values, further analysis is needed to interpret the significance of individual independent variables based on their coefficients.

```

96 # R.2.3.4: Logistic Regression
97
98 # Head subset data
99 head(subset_data$PC_WBP, 1)
100 head(subset_data$PC_WBP, 1)
101 head(subset_data$age, 1)
102 head(subset_data$heats_at_home, 1)
103 head(subset_data$possession_ratio_card, 1)
104 head(subset_data$education, 1)
105 head(subset_data$footcatal_q, 1)
106
107 # The model
108
109 ## R.2.3.4: Logistic Regression
110
111 # Coefficients:
112 (Intercept) 162.4576    75.1886    4.481 1.98e-08 ***
113 roma_score   3.8890     0.3177    12.246 8.95e-09 ***
114
115 Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
116
117 Residual standard error: 363.2 on 10 degrees of freedom
118 Multiple R-squared:  0.2155,    Adjusted R-squared:  0.1403
119 F-statistic: 18.17 on 2 and 10 DF, p-value: 0.000001
120
121 # head(subset_data$PC_WBP, 1)
122 [1] 2000.12
123 # head(subset_data$PC_WBP, 1)
124 [1] 2000.75
125 # head(subset_data$age, 1)
126 [1] 28
127 # head(subset_data$heats_at_home, 1)
128 [1] 80
129 # head(subset_data$possession_ratio_card, 1)
130 [1] 0
131 # head(subset_data$education, 1)
132 [1] 0
133 # head(subset_data$footcatal_q, 1)
134 [1] 0.13069
135

```

Interpretation :

indicates a statistically significant linear regression model, explaining around 39% of the dependent variable's variance. There's no evidence of multicollinearity among the independent variables based on the VIF values.

(ii) Linear Regression Analysis and Regression Diagnostics IPL

Python

```
In [26]: # Subset data to state assigned
subset_data = data[data['state_1'] == 'MIZ'][['foodtotal_v', 'hhdsz', 'Regular_salary_earner', 'MPCE_MRP', 'MPCE_URP', 'Possess_ration_card', 'Education', 'No_of_Meals_per_day']]
print(subset_data)
```

	foodtotal_v	hhdsz	Regular_salary_earner	MPCE_MRP	MPCE_URP	Possess_ration_card	Education	No_of_Meals_per_day
14581	968.718500	4	2.0	2925.13	3449.75	1.0	8.0	NaN
14582	1039.043333	3	2.0	2854.86	3621.80	1.0	8.0	2.0
14583	766.020714	7	2.0	2055.04	2026.00	1.0	8.0	2.0
14584	744.270000	2	1.0	2658.94	2562.50	1.0	7.0	2.0
14585	900.351667	3	2.0	1993.71	1943.67	1.0	7.0	2.0
...
47552	450.515833	6	2.0	845.74	824.50	1.0	5.0	2.0
47553	542.277000	5	2.0	1011.32	882.20	1.0	7.0	2.0
47554	448.071250	4	2.0	1008.17	1023.80	1.0	7.0	2.0
47555	468.479000	5	2.0	943.10	847.40	1.0	5.0	2.0
47556	538.958750	4	2.0	1062.15	981.50	1.0	6.0	2.0

[1536 rows x 8 columns]

```
In [23]: # Fit the regression model
X = subset_data[['hhdsz', 'Regular_salary_earner', 'MPCE_MRP', 'MPCE_URP', 'Possess_ration_card', 'Education', 'No_of_Meals_per_day']]
X = sm.add_constant(X) # Adds a constant term to the predictor
y = subset_data['foodtotal_v']

model = sm.OLS(y, X).fit()

# Print the regression results
print(model.summary())
```

OLS Regression Results

		R-squared:	0.503
Dep. Variable:	foodtotal_v	Adj. R-squared:	0.502
Model:	OLS	F-statistic:	1302.
Method:	Least Squares	Prob (F-statistic):	0.00
Date:	Sun, 23 Jun 2024	Log-likelihood:	-61381.
Time:	21:16:31	AIC:	1.228e+05
No. Observations:	9015	BIC:	1.228e+05
Df Residuals:	9007		
Df Model:	7		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	361.0196	23.716	15.223	0.000	314.531	407.509
hhdsz	-12.9630	0.860	-15.074	0.000	-14.649	-11.277
Regular_salary_earner	-14.6088	6.197	-2.357	0.018	-26.756	-2.462
MPCE_MRP	0.0728	0.002	34.081	0.000	0.069	0.077
MPCE_URP	0.0592	0.002	30.731	0.000	0.055	0.063
Possess_ration_card	-48.0845	5.077	-8.181	0.000	-59.606	-36.563
Education	7.6343	0.638	11.965	0.000	6.384	8.885
No_of_Meals_per_day	49.8726	0.234	6.057	0.000	33.732	66.013

Dennis: 3368.303 Durbin-Watson: 1.688
Prob(Dennis): 0.000 Jarque-Bera (JB): 1256929.686
Skew: -0.422 Prob(JB): 0.00
Kurtosis: 60.840 Cond. No. 3.22e+04

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 3.22e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Interpretation :

results from a statistical analysis examining how multiple factors (predictor variables) influence a single outcome (response variable). The analysis suggests a statistically significant relationship, but only explains around half of the outcome's variation. There might be issues with the data's normality or how the factors interact with each other.

```
In [24]: # multicollinearity using Variance Inflation Factor (VIF)
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.columns))]
print(vif_data) # VIF Value more than 5 is problematic
```

	feature	VIF
0	const	105.477846
1	hhdss	1.098855
2	Regular_salary_earner	1.138218
3	NPCE_MRP	2.068354
4	NPCE_URP	1.968635
5	Possess_ration_card	1.048881
6	Education	1.238296
7	No_of_Heals_per_day	1.004672

```
In [25]: # Extract the coefficients from the model
coefficients = model.params

# Construct the equation
equation = f"y = {coefficients[0]:.2f}"
for i in range(1, len(coefficients)):
    equation += f" + {coefficients[i]:.6f}*x{i}"

# Print the equation
print(equation)

y = 361.02 + -12.963010*x1 + -14.608830*x2 + 0.072781*x3 + 0.059190*x4 + -48.084503*x5 + 7.634267*x6 + 49.872590*x7
```

Interpretation :

This linear regression analysis suggests several factors influence the outcome, but there's room for improvement. The model is statistically significant, though it only explains about half the outcome's variation. It's possible the data isn't perfectly normal or the factors themselves interact in unexpected ways.

R

```
Call:
lm(formula = y_train_runs ~ runs_scored, data = data.frame(runs_scored = X_train_runs$runs_scored,
y_train_runs))

Residuals:
    Min       1Q   Median       3Q      Max
-851.2  -316.8  -127.1   346.3  1053.5

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  332.8328    75.5888   4.403 5.08e-05 ***
runs_scored    1.3690     0.3177   4.310 6.97e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 463.2 on 54 degrees of freedom
Multiple R-squared:  0.2559,    Adjusted R-squared:  0.2421
F-statistic: 18.57 on 1 and 54 DF,  p-value: 6.967e-05
```

statistical analysis was performed to see how runs scored affects a dependent variable (y_train_runs). The results show a positive correlation, meaning more runs lead to higher values in y_train_runs. The model itself is statistically significant, but only explains a quarter of the variation in y_train_runs.

```

96
97 # create a linear regression model for wickets
98 model_wickets <- lm(y_train_wickets ~ wicket_confirmation, data = data.frame(wicket_confirmation = X_train_wickets$wicket_confirmation,
99 y_train_wickets = y_train_wickets))
100 print(summary_wickets)
101
102 # evaluate the model for runs
103 y_pred_runs <- predict(model_runs, newdata = data.frame(runs_scored = X_test_runs$runs_scored))
104 r2_runs <- cor(y_test_runs, y_pred_runs)^2
105 print(paste("R-squared for runs: ", r2_runs))
106
107 # Evaluate the model for wickets
108 y_pred_wickets <- predict(model_wickets, newdata = data.frame(wicket_confirmation = X_test_wickets$wicket_confirmation))
109 r2_wickets <- cor(y_test_wickets, y_pred_wickets)^2
110 print(paste("R-squared for wickets: ", r2_wickets))

```

10022 (Top Level):

Console Terminal Background Jobs

```

R 4.3.2 C:\Users\Chand\Downloads\Assignment 1b\
> y_test_wickets <- y_wickets[-trainindex_wickets]
> # Create a linear regression model for wickets
> model_wickets <- lm(y_train_wickets ~ wicket_confirmation, data = data.frame(wicket_confirmation = X_train_wickets$wicket_confirmation,
y_train_wickets = y_train_wickets))
> summary_wickets <- summary(model_wickets)
> print(summary_wickets)

call:
lm(formula = y_train_wickets ~ wicket_confirmation, data = data.frame(wicket_confirmation = X_train_wickets$wicket_confirmation,
y_train_wickets))

Residuals:
    Min       1Q   Median       3Q      Max
-543.7  -215.3  -142.3   207.7   856.7

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)    139.215     87.222   1.596   0.1185
wicket_confirmation  26.530      7.326   3.525   0.0011 **
---
signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 363.2 on 39 degrees of freedom
Multiple R-squared:  0.2416,    Adjusted R-squared:  0.2222
F-statistic: 12.43 on 1 and 39 DF,  p-value: 0.001099

```

- A statistically significant positive correlation between `wicket_confirmation` and `y_train_wickets`. This means that higher values in `wicket_confirmation` are associated with higher values in `y_train_wickets`.
- The model explains about 24% of the variance in `y_train_wickets`.

```

> # Evaluate the model for runs
> y_pred_runs <- predict(model_runs, newdata = data.frame(runs_scored = X_test_runs$runs_scored))
> r2_runs <- cor(y_test_runs, y_pred_runs)^2
> print(paste("R-squared for runs: ", r2_runs))
[1] "R-squared for runs: 0.190229134838644"
> # Evaluate the model for wickets
> y_pred_wickets <- predict(model_wickets, newdata = data.frame(wicket_confirmation = X_test_wickets$wicket_confirmation))
> r2_wickets <- cor(y_test_wickets, y_pred_wickets)^2
> print(paste("R-squared for wickets: ", r2_wickets))
[1] "R-squared for wickets: 0.155208492981248"
>

```

the code is calculating the R-squared value to see how well the linear model fits the data for predicting test runs. An R-squared value of 0.19 indicates that the model explains only about 19% of the variance in the test runs. And 0.15, 15% of variance for wickets

```
> # Print the equation
> print(equation)
[1] "y = 139.22 + 26.52955*x1"
> |
```

The equation is:

$$y = 139.22 + 26.53x$$

In this equation:

- y represents the predicted value
- x represents the independent variable

Recommendations :

Address Normality Issues: Since the diagnostic tests highlighted potential normality issues such as high Jarque-Bera and Kurtosis values, it is recommended to explore transformations or robust regression techniques to address the non-normality in the data .

Manage Multicollinearity: Given the indication of multicollinearity in the model (high condition number), it is advisable to consider techniques like ridge regression or principal component analysis to mitigate the effects of multicollinearity and improve the stability of the regression coefficients .

Further Model Refinement: To enhance the predictive power of the model, additional variables or interaction terms could be considered based on domain knowledge or further data exploration. This may help in capturing more nuances in the relationship between predictors and the response variable .

Validation and Sensitivity Analysis: It is essential to validate the model on independent datasets or through cross-validation techniques to ensure its generalizability. Sensitivity analysis can also be conducted to assess the robustness of the model to different assumptions or variations in the data .

Continuous Monitoring: Regular monitoring of the model performance and updating it with new data can help in maintaining its relevance and accuracy over time. This iterative process can lead to continuous improvement in predictions and insights derived from the model.

Regression Analysis in NSSO :

5. Codes :

R :

```
# Set the working directory and verify it

#NSSO

#Dplyr

library(dplyr)

setwd('C:\\Users\\Chand\\Downloads\\Assignment1')

getwd()


# Load the dataset

data <- read.csv("NSSO68.csv")

unique(data$state_1)


# Subset data to state assigned

subset_data <- data %>%

  filter(state_1 == 'MIZ') %>%

  select(foodtotal_q, MPCE_MRP,
MPCE_URP, Age, Meals_At_Home, Possess_ration_card, Education, No_of_Meals_per_day)

print(subset_data)


sum(is.na(subset_data$MPCE_MRP))

sum(is.na(subset_data$MPCE_URP))

sum(is.na(subset_data$Age))

sum(is.na(subset_data$Possess_ration_card))

sum(is.na(data$Education))


impute_with_mean <- function(data, columns) {

  data %>%

    mutate(across(all_of(columns), ~ ifelse(is.na(.), mean(., na.rm = TRUE), .)))

}
```

```

# Columns to impute
columns_to_impute <- c("Education")

# Impute missing values with mean
data <- impute_with_mean(data, columns_to_impute)
sum(is.na(data$Education))

# Fit the regression model
model <- lm(foodtotal_q~
MPCE_MRP+MPCE_URP+Age+Meals_At_Home+Possess_ration_card+Education, data =
subset_data)

# Print the regression results
print(summary(model))

library(car)
# Check for multicollinearity using Variance Inflation Factor (VIF)
vif(model) # VIF Value more than 8 its problematic

# Extract the coefficients from the model
coefficients <- coef(model)

# Construct the equation
equation <- paste0("y = ", round(coefficients[1], 2))
for (i in 2:length(coefficients)) {
  equation <- paste0(equation, " + ", round(coefficients[i], 6), "*x", i-1)
}
# Print the equation
print(equation)

```

```
head(subset_data$MPCE_MRP,1)
```

```
head(subset_data$MPCE_URP,1)
```

```
head(subset_data$Age,1)
```

```
head(subset_data$Meals_At_Home,1)
```

```
head(subset_data$Possess_ration_card,1)
```

```
head(subset_data$Education,1)
```

```
head(subset_data$foodtotal_q,1)
```

Python :

```
In [1]: # Import pandas as pd, numpy as np

In [2]: # Import os
os.chdir('C:\\Users\\Chand\\Downloads\\Assignment 1b\\')

In [3]: # df IPL = pd.read_csv("IPL_ball_by_ball_updated_till_2024.csv",low_memory=False)
salary = pd.read_excel("IPL SALARIES 2024.xlsx")

In [4]: # df IPL.columns
Out[4]: Index(['Match Id', 'Date', 'Season', 'Batting team', 'Bowling team',
              'Innings No', 'Ball No', 'Bowler', 'Striker', 'Non Striker',
              'runs_scored', 'extras', 'type of extras', 'score', 'score/wicket',
              'wicket_confirmation', 'wicket_type', 'fielders_involved',
              'Player Out'],
              dtype='object')

In [5]: # grouped_data = df IPL.groupby(['Season', 'Innings No', 'Striker', 'Bowler']).agg(['runs_scored': sum, 'wicket_confirmation': sum])

In [6]: # grouped_data
Out[6]:
```

	Season	Innings No	Striker	Bowler	runs_scored	wicket_confirmation
0	2007/08	1	A Chopra	DF Vijaykumar	1	0
1	2007/08	1	A Chopra	DW Steyn	1	1
2	2007/08	1	A Chopra	DD McGrath	2	0
3	2007/08	1	A Chopra	PJ Sargeant	6	1
4	2007/08	1	A Chopra	RP Singh	8	0
...
48781	2024	2	YBK Jaswal	R/W Topley	0	1
48782	2024	2	YBK Jaswal	SM Curran	8	0
48783	2024	2	YBK Jaswal	Tiel Verna	5	0
48784	2024	2	YBK Jaswal	YG Arora	10	1
48785	2024	2	YBK Jaswal	Yash Thakur	5	0

48786 rows x 6 columns

```
In [7]: # total_runs_each_year = grouped_data.groupby(['Season', 'Striker'])['runs_scored'].sum().reset_index()
total_wicket_each_year = grouped_data.groupby(['Season', 'Bowler'])['wicket_confirmation'].sum().reset_index()
```

```
In [8]: # total_runs_each_year
Out[8]:
```

	Season	Striker	runs_scored
0	2007/08	A Chopra	42
1	2007/08	A Kumble	13
2	2007/08	A Mishra	37
3	2007/08	A Mukund	0
4	2007/08	A Nehra	3
...
2583	2024	Vijaykumar Vysakh	1
2584	2024	WG Jacks	176
2585	2024	WR Saha	135
2586	2024	Washington Sundar	0
2587	2024	YBK Jaswal	248

2598 rows x 3 columns

```
In [9]: # pip install python-Levenshtein

In [10]: # pip install python-Levenshtein
Requirement already satisfied: python-Levenshtein in c:\users\chand\anaconda3\lib\site-packages (0.25.1)
Requirement already satisfied: Levenshtein==0.25.1 in c:\users\chand\anaconda3\lib\site-packages (from python-Levenshtein) (0.25.1)
Requirement already satisfied: rapidfuzz<4.0.0,>3.0.0 in c:\users\chand\anaconda3\lib\site-packages (from Levenshtein==0.25.1->python-Levenshtein) (3.0.3)
Note: you may need to restart the kernel to use updated packages.

In [11]: # from fuzzywuzzy import process

# Convert to DataFrame
df_salary = salary.copy()
df_runs = total_runs_each_year.copy()

# Function to match names
def match_names(name, names_list):
    match, score = process.extractOne(name, names_list)
    return match if score >= 80 else None # Use a threshold score of 80

# Create a new column in df_salary with matched names from df_runs
df_salary['Matched_Player'] = df_salary['Player'].apply(lambda x: match_names(x, df_runs['Striker'].tolist()))

# Merge the DataFrames on the matched names
```

```
In [10]: ! pip install python-Levenshtein

Requirement already satisfied: python-Levenshtein in c:\users\chand\anaconda3\lib\site-packages (0.25.1)
Requirement already satisfied: Levenshtein==0.25.1 in c:\users\chand\anaconda3\lib\site-packages (from python-Levenshtein) (0.25.1)
Requirement already satisfied: rapidfuzz<4.0.0,>=3.8.0 in c:\users\chand\anaconda3\lib\site-packages (from Levenshtein==0.25.1->python-Levenshtein) (3.9.3)
Note: you may need to restart the kernel to use updated packages.
```

```
In [11]: ! from fuzzywuzzy import process

# Convert to DataFrame
df_salary = salary.copy()
df_runs = total_runs_each_year.copy()

# Function to match names
def match_names(name, names_list):
    match, score = process.extractOne(name, names_list)
    return match if score >= 80 else None # Use a threshold score of 80

# Create a new column in df_salary with matched names from df_runs
df_salary['Matched_Player'] = df_salary['Player'].apply(lambda x: match_names(x, df_runs['Striker'].tolist()))

# Merge the DataFrames on the matched names
df_merged = pd.merge(df_salary, df_runs, left_on='Matched_Player', right_on='Striker')
```

```
In [14]: ! df_original = df_merged.copy()
```

```
In [15]: ! #subsets data for last three years
df_merged = df_merged.loc[df_merged['Season'].isin(['2021', '2022', '2023'])]
```

```
In [16]: ! df_merged.Season.unique()
```

```
Out[16]: array(['2023', '2022', '2021'], dtype=object)
```

```
In [17]: ! df_merged.head()
```

```
Out[17]:
```

	Player	Salary	Rs	international	iconic	Matched_Player	Season	Striker	runs_scored
0	Abhishek Paree	20 lakh	20	0	NaN	Abhishek Paree	2023	Abhishek Paree	33
3	Anrich Nortje	6.5 crore	650	1	NaN	A Nortje	2022	A Nortje	1
4	Anrich Nortje	6.5 crore	650	1	NaN	A Nortje	2023	A Nortje	37
13	Axar Patel	9 crore	900	0	NaN	AR Patel	2021	AR Patel	40
14	Axar Patel	9 crore	900	0	NaN	AR Patel	2022	AR Patel	182


```
In [18]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
```

```
In [19]: import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_percentage_error
X = df_merged[['runs_scored']] # Independent variable(s)
y = df_merged['Rs'] # Dependent variable
# Split the data into training and test sets (80% for training, 20% for testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Create a LinearRegression model
model = LinearRegression()
# Fit the model on the training data
model.fit(X_train, y_train)
```

Out[19]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [20]: X.head()
```

```
Out[20]:
```

	runs_scored
0	33
3	1
4	37
13	40
14	182

```
In [21]: import pandas as pd
from sklearn.model_selection import train_test_split
import statsmodels.api as sm

# Assuming df_merged is already defined and contains the necessary columns
X = df_merged[['runs_scored']] # Independent variable(s)
y = df_merged['Rs'] # Dependent variable

# Split the data into training and test sets (80% for training, 20% for testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Add a constant to the model (intercept)
X_train_sm = sm.add_constant(X_train)

# Create a statsmodels OLS regression model
model = sm.OLS(y_train, X_train_sm).fit()
```

```
In [25]: import pandas as pd
from sklearn.model_selection import train_test_split
import statsmodels.api as sm

# Assuming df_merged is already defined and contains the necessary columns
X = df_merged[['wicket_confirmation']] # Independent variable(s)
y = df_merged['Rs'] # Dependent variable

# Split the data into training and test sets (80% for training, 20% for testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Add a constant to the model (intercept)
X_train_sm = sm.add_constant(X_train)

# Create a statsmodels OLS regression model
model = sm.OLS(y_train, X_train_sm).fit()

# Get the summary of the model
summary = model.summary()
print(summary)
```

```

OLS Regression Results
=====
Dep. Variable:          Rs      R-squared:          0.074
Model:                  OLS      Adj. R-squared:       0.054
Method:                 Least Squares      F-statistic:         3.688
Date:                   Sun, 23 Jun 2024      Prob (F-statistic):    0.0610
Time:                   13:43:25      Log-Likelihood:       -360.96
No. Observations:       48      AIC:                  725.9
Df Residuals:           46      BIC:                  729.7
Df Model:                1
Covariance Type:        nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const          396.6881      91.270      4.346      0.000      212.971      580.405
wicket_confirmation      17.6635      9.198      1.928      0.061      -0.851      36.179
=====
Omnibus:                6.984      Durbin-Watson:         2.451
Prob(Omnibus):           0.030      Jarque-Bera (JB):        6.309
Skew:                   0.877      Prob(JB):               0.0427
Kurtosis:                3.274      Cond. No.               13.8
=====
```

```
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

Regression Analysis in IPL:

R:

Load necessary libraries

install.packages("stringdist")

library(readr)

library(readxl)

library(dplyr)

library(stringdist)

Change the directory to where the datasets are stored

setwd('C:\\Users\\Chand\\Downloads\\Assignment 1b')

Load the datasets

```
df_ipl <- read_csv("IPL_ball_by_ball_updated till 2024.csv")
```

```
salary <- read_excel("IPL SALARIES 2024.xlsx")
```

```
# Group and aggregate the performance metrics
```

```
grouped_data <- df_ipl %>%
```

```
  group_by(Season, `Innings No`, Striker, Bowler) %>%
```

```
  summarise(
```

```
    runs_scored = sum(runs_scored, na.rm = TRUE),
```

```
    wicket_confirmation = sum(wicket_confirmation, na.rm = TRUE)
```

```
  ) %>%
```

```
  ungroup()
```

```
# Calculate total runs and wickets each year
```

```
total_runs_each_year <- grouped_data %>%
```

```
  group_by(Season, Striker) %>%
```

```
  summarise(runs_scored = sum(runs_scored, na.rm = TRUE)) %>%
```

```
  ungroup()
```

```
total_wicket_each_year <- grouped_data %>%
```

```
  group_by(Season, Bowler) %>%
```

```
  summarise(wicket_confirmation = sum(wicket_confirmation, na.rm = TRUE)) %>%
```

```
  ungroup()
```

```
# Function to match names
```

```
match_names <- function(name, names_list) {
```

```
  match <- amatch(name, names_list, maxDist = 0.2)
```

```
  if (!is.na(match)) {
```

```
    return(names_list[match])
```

```
  } else {
```

```
    return(NA)
```

```
}  
}
```

```
# Matching names for runs
```

```
df_salary_runs <- salary
```

```
df_runs <- total_runs_each_year
```

```
df_salary_runs$Matched_Player <- sapply(df_salary_runs$Player, function(x)  
match_names(x, df_runs$Striker))
```

```
# Merge the DataFrames for runs
```

```
df_merged_runs <- merge(df_salary_runs, df_runs, by.x = "Matched_Player", by.y =  
"Striker")
```

```
# Subset data for the last three years
```

```
df_merged_runs <- df_merged_runs %>% filter(Season %in% c("2021", "2022", "2023"))
```

```
# Perform regression analysis for runs
```

```
X_runs <- df_merged_runs %>% select(runs_scored)
```

```
y_runs <- df_merged_runs$Rs
```

```
# Split the data into training and test sets (80% for training, 20% for testing)
```

```
set.seed(42)
```

```
trainIndex_runs <- sample(seq_len(nrow(X_runs)), size = 0.8 * nrow(X_runs))
```

```
X_train_runs <- X_runs[trainIndex_runs, , drop = FALSE]
```

```
X_test_runs <- X_runs[-trainIndex_runs, , drop = FALSE]
```

```
y_train_runs <- y_runs[trainIndex_runs]
```

```
y_test_runs <- y_runs[-trainIndex_runs]
```

```
# Create a linear regression model for runs
```

```
model_runs <- lm(y_train_runs ~ runs_scored, data = data.frame(runs_scored =  
X_train_runs$runs_scored, y_train_runs))
```

```

summary_runs <- summary(model_runs)
print(summary_runs)

# Matching names for wickets
df_salary_wickets <- salary
df_wickets <- total_wicket_each_year
df_salary_wickets$Matched_Player <- sapply(df_salary_wickets$Player, function(x)
match_names(x, df_wickets$Bowler))

# Merge the DataFrames for wickets
df_merged_wickets <- merge(df_salary_wickets, df_wickets, by.x = "Matched_Player", by.y
= "Bowler")

# Subset data for the last three years
df_merged_wickets <- df_merged_wickets %>% filter(Season %in% c("2021", "2022",
"2023"))

# Perform regression analysis for wickets
X_wickets <- df_merged_wickets %>% select(wicket_confirmation)
y_wickets <- df_merged_wickets$Rs

# Split the data into training and test sets (80% for training, 20% for testing)
trainIndex_wickets <- sample(seq_len(nrow(X_wickets)), size = 0.8 * nrow(X_wickets))
X_train_wickets <- X_wickets[trainIndex_wickets, , drop = FALSE]
X_test_wickets <- X_wickets[-trainIndex_wickets, , drop = FALSE]
y_train_wickets <- y_wickets[trainIndex_wickets]
y_test_wickets <- y_wickets[-trainIndex_wickets]

# Create a linear regression model for wickets
model_wickets <- lm(y_train_wickets ~ wicket_confirmation, data =
data.frame(wicket_confirmation = X_train_wickets$wicket_confirmation, y_train_wickets))
summary_wickets <- summary(model_wickets)

```

```
print(summary_wickets)
```

```
# Evaluate the model for runs
```

```
y_pred_runs <- predict(model_runs, newdata = data.frame(runs_scored =  
X_test_runs$runs_scored))
```

```
r2_runs <- cor(y_test_runs, y_pred_runs)^2
```

```
print(paste("R-squared for runs: ", r2_runs))
```

```
# Evaluate the model for wickets
```

```
y_pred_wickets <- predict(model_wickets, newdata = data.frame(wicket_confirmation =  
X_test_wickets$wicket_confirmation))
```

```
r2_wickets <- cor(y_test_wickets, y_pred_wickets)^2
```

```
print(paste("R-squared for wickets: ", r2_wickets))
```

Python :

```
In [1]: # Import statsmodels.api as sm  
from statsmodels.stats.outliers_influence import variance_inflation_factor  
from sklearn.impute import SimpleImputer  
import os  
import pandas as pd
```

```
In [2]: # Set working directory  
os.chdir('C:\\Users\\Chand\\Downloads\\Assignment1')  
print(os.getcwd())  
C:\\Users\\Chand\\Downloads\\Assignment1
```

```
In [3]: # Load the dataset  
data = pd.read_csv("N55008.csv")  
C:\\Users\\Chand\\AppData\\Local\\Temp\\ipykernel_15052\\95200774.py:2: DtypeWarning: Columns (1) have mixed types. Specify dtype option on import or set low_memory=False.  
data = pd.read_csv("N55008.csv")
```

```
In [26]: # Subset data to state assigned  
subset_data = data[data["state_1"] == "MIZ"]  
print(subset_data)
```

	foodtotal_v	hhdsz	Regular_salary_earner	MPCE_HRP	MPCE_URP
14581	968.718500	4	2.0	2025.13	3440.75
14582	1039.043333	3	2.0	2054.06	3021.00
14583	766.020714	7	2.0	2055.04	2026.00
14584	744.270000	2	1.0	2658.94	2562.50
14585	900.351667	3	2.0	1993.71	1943.67
...
47552	450.515000	6	1.0	845.74	824.50
47553	542.277000	5	2.0	1011.32	882.20
47554	448.071250	4	2.0	1019.07	1023.00
47555	468.479000	5	2.0	943.10	847.40
47556	538.958750	4	2.0	1045.80	981.50

	Possess_ration_card	Education	No_of_Meals_per_day
14581	1.0	8.0	NaN
14582	1.0	6.0	2.0
14583	1.0	8.0	2.0
14584	1.0	7.0	2.0
14585	1.0	7.0	2.0
...
47552	1.0	5.0	2.0
47553	1.0	7.0	2.0
47554	1.0	7.0	2.0
47555	1.0	7.0	2.0
47556	1.0	7.0	2.0

```
In [22]: # Check for missing values
print(subset_data['hhdsz'].isna().sum())
print(subset_data['Regular_salary_earner'].isna().sum())
print(subset_data['MPCE_HRP'].isna().sum())
print(subset_data['MPCE_URP'].isna().sum())
print(subset_data['Possess_ration_card'].isna().sum())
print(subset_data['Education'].isna().sum())
print(subset_data['No_of_Meals_per_day'].isna().sum())

# Impute missing values with mean
imputer = SimpleImputer(strategy='mean')
subset_data['Possess_ration_card'] = imputer.fit_transform(subset_data[['Possess_ration_card']])

print("Possess_ration_card:")
print(subset_data['Possess_ration_card'].isna().sum())

0
0
0
0
2
0
0
Possess_ration_card:
0
```

```
In [23]: # Fit the regression model
X = subset_data[['hhdsz', 'Regular_salary_earner', 'MPCE_HRP', 'MPCE_URP', 'Possess_ration_card', 'Education', 'No_of_Meals_per_day']]
X = sm.add_constant(X) # Adds a constant term to the predictor
y = subset_data['foodtotal_v']

model = sm.OLS(y, X).fit()

# Print the regression results
print(model.summary())
```

```
OLS Regression Results
```

Dep. Variable:	foodtotal_v	R-squared:	0.503
Model:	OLS	Adj. R-squared:	0.502
Method:	Least Squares	F-statistic:	1302.
Date:	Sun, 23 Jun 2024	Prob (F-statistic):	0.00
Time:	21:16:31	Log-Likelihood:	-61381.
No. Observations:	9015	AIC:	1.228e+05
Df Residuals:	9007	BIC:	1.228e+05
Df Model:	7		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	361.02	12.903010	27.978	0.000	335.12	386.92
hhdsz	-14.608830	0.872781	-16.739	0.000	-16.342	-12.876
Regular_salary_earner	0.872781	0.850190	1.027	0.305	-0.804	2.549
MPCE_HRP	-48.084583	7.634267	-6.300	0.000	-62.994	-33.174
MPCE_URP	49.872596	12.903010	3.865	0.000	24.074	75.679
Possess_ration_card	1.230296	1.004672	1.224	0.222	-0.748	3.209
Education	1.004672	1.230296	0.817	0.415	-1.438	3.450
No_of_Meals_per_day	1.004672	1.004672	1.000	0.317	-0.995	3.007

[4] THE COLLINUM NUMBER IS LARGE, 2.22E+05. THIS MIGHT INDICATE THAT THERE ARE STRONG MULTICOLLINEARITY OR OTHER NUMERICAL PROBLEMS.

```
In [24]: # Multicollinearity using Variance Inflation Factor (VIF)
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.columns))]
print(vif_data) # VIF Value more than 8 is problematic
```

	feature	VIF
0	const	105.477846
1	hhdsz	1.008855
2	Regular_salary_earner	1.138218
3	MPCE_HRP	2.008354
4	MPCE_URP	1.968635
5	Possess_ration_card	1.048881
6	Education	1.230296
7	No_of_Meals_per_day	1.004672

```
In [25]: # Extract the coefficients from the model
coefficients = model.params

# Construct the equation
equation = f'y = {coefficients[0]:.2f}'
for i in range(1, len(coefficients)):
    equation += f' + {coefficients[i]:.0f}*x{i}'

# Print the equation
print(equation)

y = 361.02 + -12.903010*x1 + -14.608830*x2 + 0.872781*x3 + 0.850190*x4 + -48.084583*x5 + 7.634267*x6 + 49.872596*x7
```

6. References :

Statistical analysis and modelling (SCMA 632)